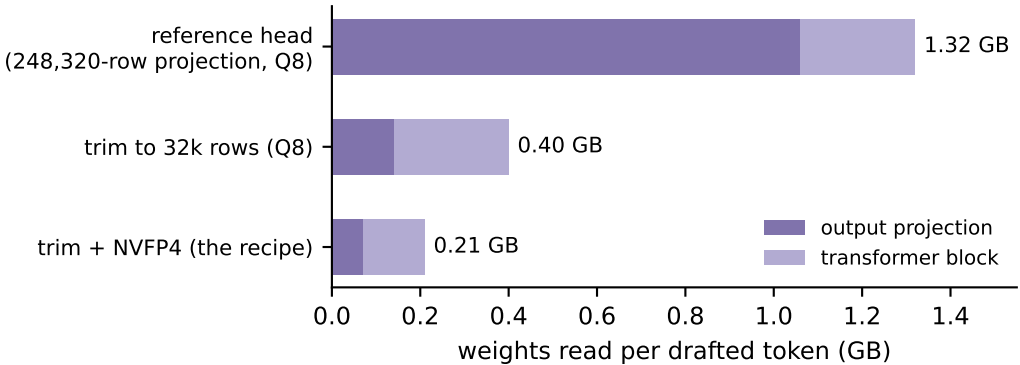
reference head
(248,320-row projection, Q8)
trim to 32k rows (Q8)
trim + NVFP4 (the recipe)
1.32 GB
0.40 GB
0.21 GB
output projection
transformer block
0.0
0.2
0.4
0.6
0.8
1.0
1.2
1.4
weights read per drafted token (GB)